

4-th Exercise - Embeddings, Vector Databases & Search

We will use FAISS and CLIP to create and search pre-trained embedding vectors.

First, we need to install required libraries:

- CLIP: For generating embeddings
- FAISS: For indexing

```
In [1]: import sys

# Uncomment as needed
# !{sys.executable} -m pip install --no-input openai-clip

# Either use CPU-only version
# !conda install --yes --prefix {sys.prefix} -c pytorch faiss-cpu=1.9.0

# Alternatively use the GPU(+CPU) version
# !conda install --yes --prefix {sys.prefix} -c pytorch -c nvidia faiss-gpu=1.9.0

# !conda install --yes --prefix {sys.prefix}
```

Load required libraries

```
In [2]: import os
# Avoid different versions of numpy make kernel collapses
os.environ['KMP_DUPLICATE_LIB_OK']=True

import torch
import faiss
import clip
import json
import requests

from PIL import Image
from io import BytesIO

import time
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from matplotlib.patches import Rectangle

from tqdm import tqdm
```

Some helper methods for testing results

```
In [3]: images_URL = "http://vision.stanford.edu/aditya86/ImageNetDogs/images/"
```

```

def load_embeddings(embeddings_path):
    df = pd.read_csv(embeddings_path, sep=';', compression='gzip', index_col=0)

    # Next, we will convert the embeddings String-column into a numpy-vector
    start_time = time.time()
    data = df["embedding"].apply(json.loads).values
    embeddings = np.zeros((data.shape[0], len(data[0])), dtype=np.float32)
    for i, d in enumerate(data):
        embeddings[i] = d

    end_time = time.time()
    print(f"Time to load {np.round(end_time-start_time,3)}s")

    return df, embeddings


# Calculate Precision@k
def precision_at_k(indices, k):
    y_true = df.iloc[indices[:,0]]["class"].values
    y_pred = df.iloc[indices.flatten()]["class"].values.reshape(indices.shape[0])
    precisions = []
    # Check how many of the retrieved items are relevant
    for i, pred in enumerate(y_pred[:, :k]):
        relevant = np.sum(y_true[i] == pred)
        precisions.append(relevant / k)
    return np.mean(precisions)


# Plot search results
def plot_results(query_file_names, indices, k):
    if indices is not None:
        for pos, ind in enumerate(indices):
            fig, ax = plt.subplots(1, k+1, figsize=(3*(k+1),4))

            if query_file_names:
                path_query = f'{query_file_names[pos]}'
                if os.path.isfile(path_query):
                    ax[0].imshow(mpimg.imread(path_query))
                    ax[0].axis('off')
                    ax[0].set_title("Query")
                elif len(query_file_names) > pos:
                    # ax[0].set_title()
                    ax[0].text(0.5, 0.5,
                               f"{query_file_names[pos]}", fontsize=20,
                               ha='center', va='center')
                    ax[0].axis('off')
            else:
                ax[0].remove()

            for i in range(0, k):
                d_train = df.iloc[ind[i]]

                url = f'{images_URL}/{d_train["dir"]}/{d_train["filename"]}.jpg'
                response = requests.get(url, stream=True)
                if response.status_code == 200:
                    img = Image.open(BytesIO(response.content))
                    ax[1+i].imshow(img)
                    ax[1+i].axis('off')
                    ax[1+i].set_title(f"{i+1}-NN: " + d_train["class"])

                xy = (d_train["xmin"], d_train["ymin"])

```

```

width = d_train["xmax"]-d_train["xmin"]
height = d_train["ymax"]-d_train["ymin"]
rect = Rectangle(xy, width, height, edgecolor='white', fill=
ax[1+i].add_patch(rect)
ax[1+i].text(xy[0], xy[1], d_train["class"], fontsize=12, co
verticalalignment='bottom', horizontalalignment='center')

else:
    ax[i+1].remove()

plt.tight_layout()
plt.show()

def plot_precision_at_k(indices, k):
    precisions = [precision_at_k(indices, i) for i in range (1,k+1)]
    for i, precision in enumerate(precisions):
        print(f"precision@{i}: {precision}")

    fig, ax = plt.subplots(figsize=(8,4))
    ax.set_title("Precision@k")
    ax.plot(np.arange(1,k+1, dtype=np.int32), precisions)
    ax.set_xlabel("k")
    ax.set_ylabel("Precision@k")
    plt.show()

```

Load Embeddings and Meta-Data

You must download the requires files first, and adapt paths:

- `queries/` : Queries needed for Tasks 1-4. Load from Moodle
- `data/vectors.csv.gz` : Embeddings needed for Tasks 1-5. Load from Moodle

```
In [4]: ### modify this path, if needed!
queries_start_path = "queries"
embeddings_path = 'data/vectors.csv.gz'

df, embeddings = load_embeddings(embeddings_path)
embedding_dim = embeddings.shape[1]
df.head()
```

Time to load 3.362s

Out[4]:

		filename	dir	class	pose	xmin	ymin	xmax	ymax
0	n02097658_26	n02097658-silky_terrier	silky_terrier	Unspecified	41	30	296	398	[0.2
1	n02097658_4869	n02097658-silky_terrier	silky_terrier	Unspecified	43	270	321	498	[-0.
2	n02097658_595	n02097658-silky_terrier	silky_terrier	Unspecified	82	7	378	355	[0.1
3	n02097658_9222	n02097658-silky_terrier	silky_terrier	Unspecified	0	12	331	498	[0.2
4	n02097658_422	n02097658-silky_terrier	silky_terrier	Unspecified	146	10	356	332	[0.0
									-0.1



Task 1: Exact Search using Flat Index

FAISS implements several metrics and indices for similarity search.

Your task is to index all embeddings using `IndexFlatL2` index first.

```
faiss.IndexFlatL2(d)
...
```

You can read more about their implementation and how to use `IndexFlatL2` on their GitHub page: <https://github.com/facebookresearch/faiss/wiki>

```
In [5]: def create_IndexFlatL2(embeddings):
    # Answer:
    # faiss.IndexFlatL2(dimension)
    index = faiss.IndexFlatL2(embeddings.shape[1])
    return index

    start_time = time.time()
    index_flat = create_IndexFlatL2(embeddings)
    end_time = time.time()
    print(f"Time to index {np.round(end_time-start_time,3)}s")
```

Time to index 0.001s

Task 2: Generate embeddings for queries using OpenAI's CLIP

You are given a set of images in the `queries` subfolder. Your task is to:

- Generate CLIP-embeddings for each image
- Search for the image embeddings using the FAISS index

CLIP [1] an embedding model by OpenAI, which is used to extract a high-dimensional vector representation of an image, which captures its semantic and perceptual features.

[1] <https://github.com/openai/CLIP>

Initialize OpenAi's CLIP Encoder for feature extraction

```
In [6]: device = "cuda" if torch.cuda.is_available() else "cpu"
print("running on: " + device)
model, preprocess = clip.load("ViT-B/32", device=device)

running on: cuda
```

Now Generate Embeddings of Queries

Use CLIP to generate for all images in the queries folder the corresponding embedding.

```
model.encode_image(image)
See: https://github.com/openai/CLIP
```

Upon generation, make sure to convert all data to float32, which is required to FAISS.

```
In [22]: def create_query_embeddings(queries_start_path):

    # Load all images from the queries subfolder
    df_data = pd.DataFrame()
    query_file_names = []
    for root, dirs, files in os.walk(queries_start_path):
        for file_name in files:
            if ".jpg" in file_name:
                query_file_names.append(f"{root}/{file_name}")
                #query_file_names.append(file_name)

    # Answer:
    query_embeddings = np.zeros((0,embeddings.shape[1]), dtype=np.float32)
    for query_file_name in query_file_names:
        image = preprocess(Image.open(query_file_name)).unsqueeze(0).to(device)
        with torch.no_grad():
            image_features = model.encode_image(image)
        query_embeddings = np.vstack((query_embeddings, image_features.cpu().numpy()))
    return query_embeddings, query_file_names

query_embeddings, query_file_names = create_query_embeddings(queries_start_path)
```

Querying: Search for top-3 relevant images

We need to define a search function below to first vectorize our query image, and then search for the vectors with the closest distance.

```
index_flat.search(query_embeddings, k)
```

In [8]: `k = 3`

```
def indexed_search(index, query_embeddings, k):
    # Search for the top-3 nearest neighbors

    # Answer:
    if not index.is_trained:
        print("Training index...")
        index.train(embeddings)

    if index.ntotal == 0:
        print("Adding embeddings to index...")
        index.add(embeddings)

    distances, indices = index.search(query_embeddings, k)

    return indices

start_time = time.time()
flat_indices = indexed_search(index_flat, query_embeddings, k)
end_time = time.time()
print(f"Time to index {np.round(end_time-start_time,3)}s")
```

Adding embeddings to index...

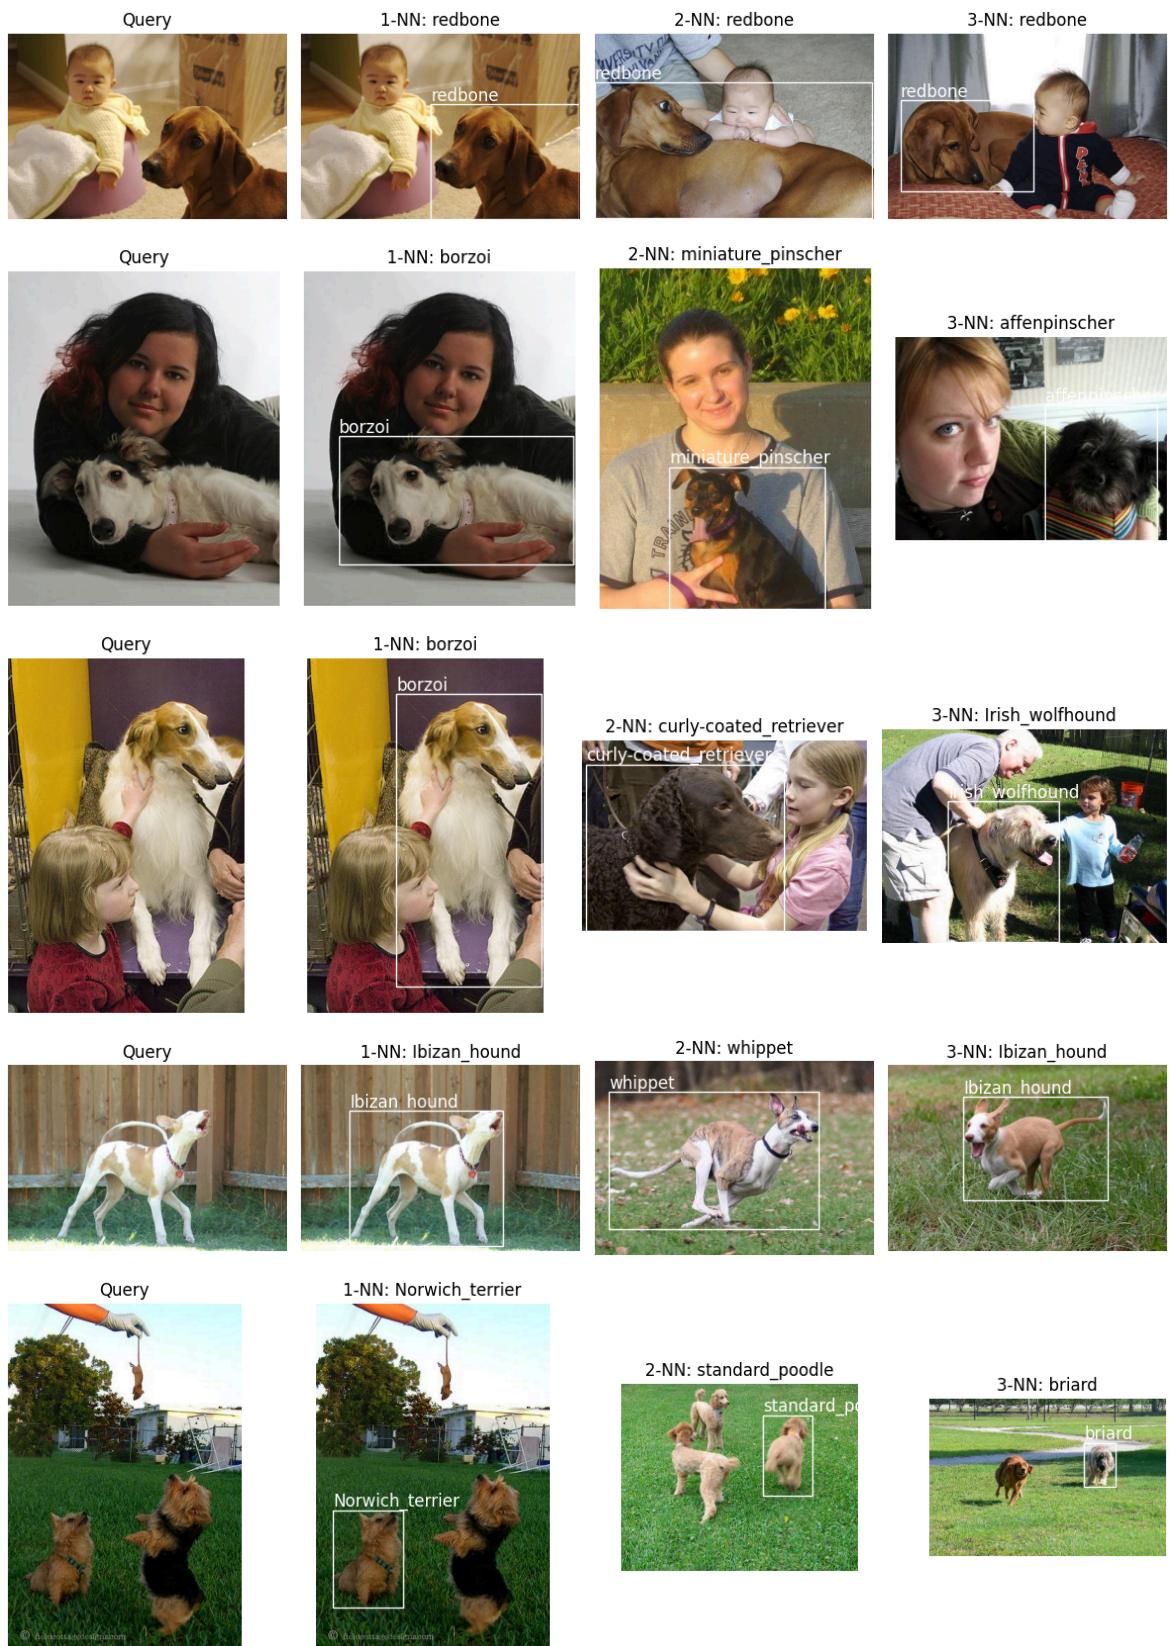
Time to index 0.021s

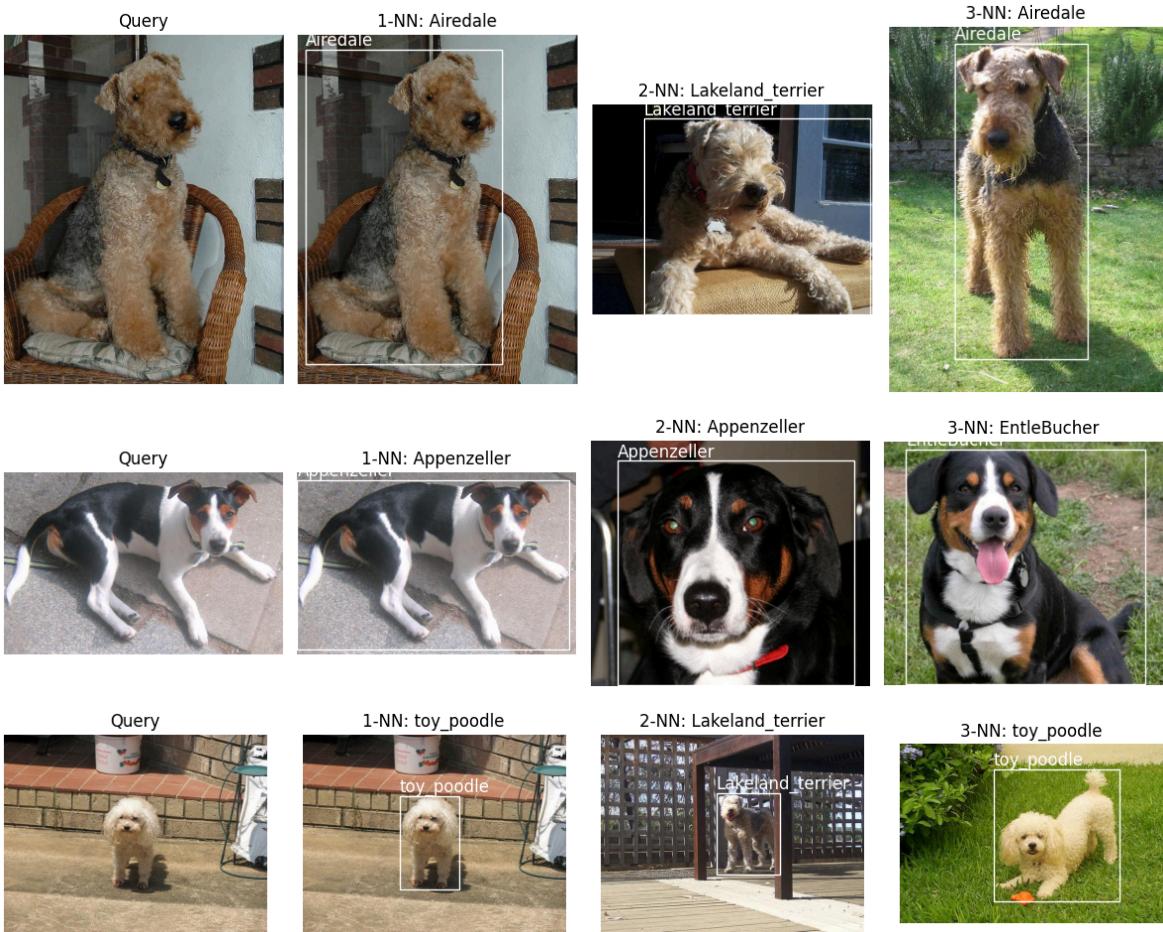
Plot results

We use the helper function, to illustrate the results

In [9]: `plot_results(query_file_names, flat_indices, k)`







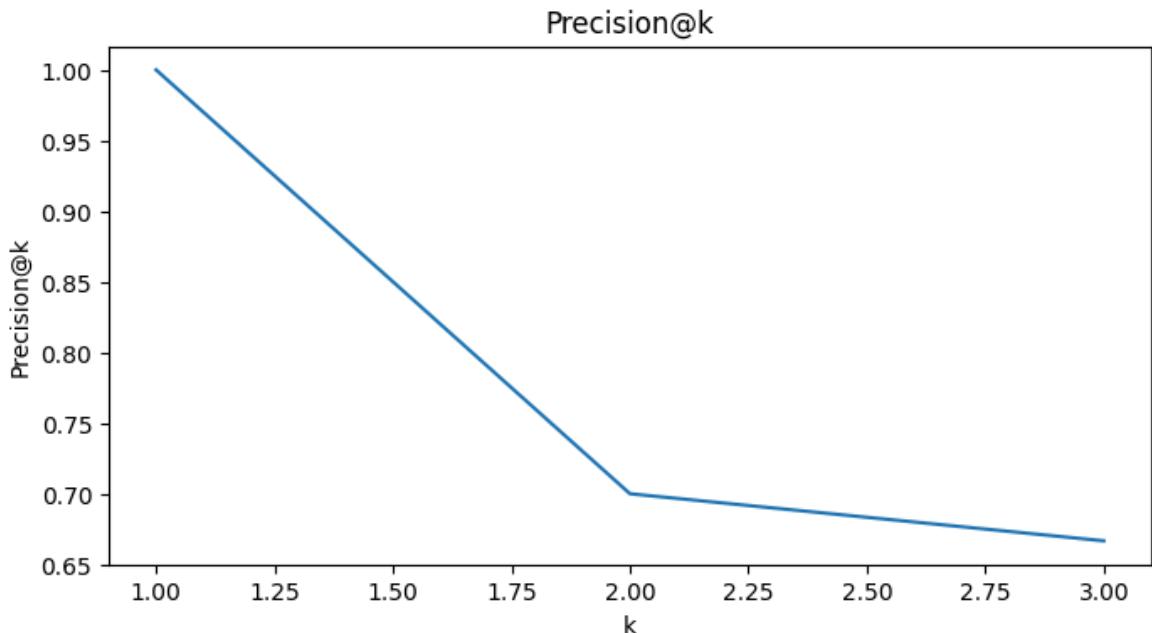
Measure precision@k

We measure the precision@k for different values of k.

```
In [10]: # A simple test to test correctness
if flat_indices is None:
    print("You must implement Task 1 and 2 first")
else:
    assert precision_at_k(flat_indices, 1) == 1.0

    plot_precision_at_k(flat_indices, k)
```

```
precision@0: 1.0
precision@1: 0.7
precision@2: 0.6666666666666667
```



Task 3: Use Hierarchical Small World Graphs

To use Hierarchical Small World Graphs in FAISS, we need to work with the HNSW (Hierarchical Navigable Small World) index that FAISS provides. HNSW is a type of small-world graph used for approximate nearest neighbor search, which is highly efficient for high-dimensional data.

```
index = faiss.IndexHNSWFlat(d, M)
hswg_index.hnsw.efConstruction = ...
hswg_index.hnsw.efSearch = ...
```

Tuning the HNSW Index:

- `M` : A higher value will give more accurate results but with increased memory usage and slower indexing times.
- `EfConstruction` : This parameter controls the number of candidate neighbors that are evaluated during the construction of the graph. A larger value can lead to better quality of the graph at the expense of longer construction time.
- `EfSearch` : This parameter controls the number of candidate neighbors considered during search. A larger value can improve accuracy but makes the search slower.

```
In [11]: k=3

def index_and_query_hswg(embeddings, query_embeddings, k):

    # Answer:
    dim = embeddings.shape[1]

    hswg_index = faiss.IndexHNSWFlat(dim, 32) # 32 is the number of neighbors to
    hswg_index.hnsw.efConstruction = 40 # 40 is the number of k-NN, when the gra
    hswg_index.hnsw.efSearch = 64 # 64 is the number of candidates to explore du

    hswg_index.add(embeddings)
```

```

distances, indices = hswg_index.search(query_embeddings, k) # used to store distances

return hswg_index, indices

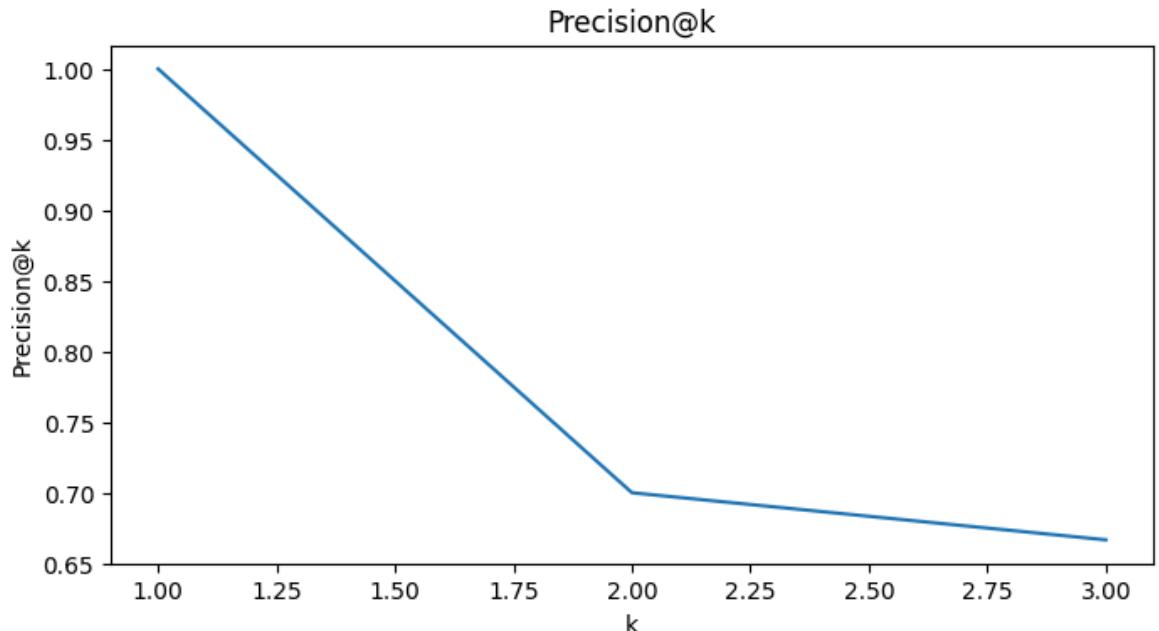
start_time = time.time()
hswg_index, hswg_indices = index_and_query_hswg(embeddings, query_embeddings, k)
end_time = time.time()
print(f"Time to index and query {np.round(end_time-start_time,3)}s")

```

Time to index and query 0.509s

```
In [12]: # A simple test to test correctness
if hswg_indices is None:
    print("You must implement Task 3 first")
else:
    assert precision_at_k(hswg_indices, 1) == 1.0
    plot_precision_at_k(hswg_indices, k)
```

precision@0: 1.0
precision@1: 0.7
precision@2: 0.6666666666666667



Task 4: Optimize Hyper-Parameterers

Test different ranges of the three input parameters

```
M_values = [8, 16, 32]
efConstruction_values = [40, 100, 200]
efSearch_values = [10, 50, 100]
metric_types = [faiss.METRIC_L2, faiss.METRIC_L1]
```

For each parameter, create a `IndexHNSWFlat`, and search for the 5-NN.

For each set of parameters, record the query-time and the precision@5. Output the best 5 sets of parameters for the best quer-times and the best precision@5.

You may measure runtimes using:

```

start_time = time.time()
...
end_time = time.time()

In [13]: # Range of parameters to try for tuning
M_values = [8, 16, 32]
efConstruction_values = [40, 100, 200]
efSearch_values = [10, 50, 100]
metric_types = [faiss.METRIC_L2, faiss.METRIC_L1]
k = 5 # Number of neighbors

def hswg_grid_search(embeddings, query_embeddings, k):
    # Answer:
    # For each set of parameters, record (a) the query-time and (b) the precision
    # Output the best 5 sets of parameters for the best quer-times and the best
    results = []
    for M in M_values:
        for efConstruction in efConstruction_values:
            for efSearch in efSearch_values:
                for metric in metric_types:
                    hswg_index = faiss.IndexHNSWFlat(embeddings.shape[1], M, metric)
                    hswg_index.hnsw.efConstruction = efConstruction
                    hswg_index.hnsw.efSearch = efSearch
                    hswg_index.reset()

                    if not hswg_index.is_trained:
                        hswg_index.train(embeddings)
                    if hswg_index.ntotal == 0:
                        print("Adding embeddings to index...")
                        hswg_index.add(embeddings)

                    start_time = time.time()
                    distances, hswg_indices = hswg_index.search(query_embeddings, k)
                    end_time = time.time()
                    query_time = end_time - start_time
                    precision = precision_at_k(hswg_indices, k)

                    results.append((M, efConstruction, efSearch, metric, query_time))

    print(f"M={M}, efConstruction={efConstruction}, efSearch={efSearch}, metric={metric}, query_time={query_time}, precision={precision}")

    # Best performance tracking
    best_params_query_time, best_params_precision_at_5 = [], []
    best_params_query_time = sorted(results, key=lambda x: x[4])[:5]
    best_params_precision_at_5 = sorted(results, key=lambda x: x[5], reverse=True)

    return best_params_query_time, best_params_precision_at_5

best_params_query_time, best_params_precision_at_5 = hswg_grid_search(embeddings)
# print ("Best Query Time", best_params_query_time)
# Best Query Time Results
print("\n== Best Query Time Results ==")
for i, result in enumerate(best_params_query_time, 1):
    print(f"{i}. Time: {result[4]:.4f}s, M={result[0]}, efc={result[1]}, efs={result[2]}, metric={result[3]}, precision={result[5]}")

# Best Precision Results
# print("\n== Best Precision Results ==")

```

```
for i, result in enumerate(best_params_precision_at_5, 1):
    print(f"{i}. Precision: {result[5]:.4f}, M={result[0]}, efC={result[1]}, efs
print ("Best Precision@5", best_params_precision_at_5)
```

```
Adding embeddings to index...
M=8, efConstruction=40, efSearch=10, metric=1: query_time=0.0, precision@5=0.4800
000000000001
Adding embeddings to index...
M=8, efConstruction=40, efSearch=10, metric=2: query_time=0.0, precision@5=0.4600
000000000001
Adding embeddings to index...
M=8, efConstruction=40, efSearch=50, metric=1: query_time=0.0, precision@5=0.5000
000000000001
Adding embeddings to index...
M=8, efConstruction=40, efSearch=50, metric=2: query_time=0.0010035037994384766,
precision@5=0.5200000000000001
Adding embeddings to index...
M=8, efConstruction=40, efSearch=100, metric=1: query_time=0.0005042552947998047,
precision@5=0.5200000000000001
Adding embeddings to index...
M=8, efConstruction=40, efSearch=100, metric=2: query_time=0.0010046958923339844,
precision@5=0.5200000000000001
Adding embeddings to index...
M=8, efConstruction=100, efSearch=10, metric=1: query_time=0.0, precision@5=0.500
000000000001
Adding embeddings to index...
M=8, efConstruction=100, efSearch=10, metric=2: query_time=0.0, precision@5=0.480
000000000001
Adding embeddings to index...
M=8, efConstruction=100, efSearch=50, metric=1: query_time=0.0, precision@5=0.520
000000000001
Adding embeddings to index...
M=8, efConstruction=100, efSearch=50, metric=2: query_time=0.0, precision@5=0.520
000000000001
Adding embeddings to index...
M=8, efConstruction=100, efSearch=100, metric=1: query_time=0.0, precision@5=0.52
000000000001
Adding embeddings to index...
M=8, efConstruction=100, efSearch=100, metric=2: query_time=0.0, precision@5=0.52
000000000001
Adding embeddings to index...
M=8, efConstruction=200, efSearch=10, metric=1: query_time=0.0005049705505371094,
precision@5=0.5200000000000001
Adding embeddings to index...
M=8, efConstruction=200, efSearch=10, metric=2: query_time=0.0, precision@5=0.480
000000000001
Adding embeddings to index...
M=8, efConstruction=200, efSearch=50, metric=1: query_time=0.0, precision@5=0.520
000000000001
Adding embeddings to index...
M=8, efConstruction=200, efSearch=50, metric=2: query_time=0.0010018348693847656,
precision@5=0.5200000000000001
Adding embeddings to index...
M=8, efConstruction=200, efSearch=100, metric=1: query_time=0.001002311706542968,
precision@5=0.5200000000000001
Adding embeddings to index...
M=8, efConstruction=200, efSearch=100, metric=2: query_time=0.001002311706542968,
precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=40, efSearch=10, metric=1: query_time=0.0, precision@5=0.500
000000000001
Adding embeddings to index...
M=16, efConstruction=40, efSearch=10, metric=2: query_time=0.0005135536193847656,
precision@5=0.5200000000000001
```

```
Adding embeddings to index...
M=16, efConstruction=40, efSearch=50, metric=1: query_time=0.0010073184967041016,
precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=40, efSearch=50, metric=2: query_time=0.0, precision@5=0.520
000000000001
Adding embeddings to index...
M=16, efConstruction=40, efSearch=100, metric=1: query_time=0.001000404357910156
2, precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=40, efSearch=100, metric=2: query_time=0.0005035400390625, p
recision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=100, efSearch=10, metric=1: query_time=0.0, precision@5=0.52
00000000000001
Adding embeddings to index...
M=16, efConstruction=100, efSearch=10, metric=2: query_time=0.001003503799438476
6, precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=100, efSearch=50, metric=1: query_time=0.001000881195068359
4, precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=100, efSearch=50, metric=2: query_time=0.0, precision@5=0.52
00000000000001
Adding embeddings to index...
M=16, efConstruction=100, efSearch=100, metric=1: query_time=0.001001358032226562
5, precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=100, efSearch=100, metric=2: query_time=0.001000165939331054
7, precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=200, efSearch=10, metric=1: query_time=0.0, precision@5=0.52
00000000000001
Adding embeddings to index...
M=16, efConstruction=200, efSearch=10, metric=2: query_time=0.0, precision@5=0.50
00000000000001
Adding embeddings to index...
M=16, efConstruction=200, efSearch=50, metric=1: query_time=0.001004219055175781
2, precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=200, efSearch=50, metric=2: query_time=0.000504970550537109
4, precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=200, efSearch=100, metric=1: query_time=0.001000165939331054
7, precision@5=0.5200000000000001
Adding embeddings to index...
M=16, efConstruction=200, efSearch=100, metric=2: query_time=0.001000881195068359
4, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=40, efSearch=10, metric=1: query_time=0.0, precision@5=0.520
000000000001
Adding embeddings to index...
M=32, efConstruction=40, efSearch=10, metric=2: query_time=0.0, precision@5=0.520
000000000001
Adding embeddings to index...
M=32, efConstruction=40, efSearch=50, metric=1: query_time=0.0, precision@5=0.520
000000000001
Adding embeddings to index...
M=32, efConstruction=40, efSearch=50, metric=2: query_time=0.0010042190551757812,
precision@5=0.5200000000000001
```

```

Adding embeddings to index...
M=32, efConstruction=40, efSearch=100, metric=1: query_time=0.001003742218017578
1, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=40, efSearch=100, metric=2: query_time=0.001004695892333984
4, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=100, efSearch=10, metric=1: query_time=0.0, precision@5=0.52
00000000000001
Adding embeddings to index...
M=32, efConstruction=100, efSearch=10, metric=2: query_time=0.0, precision@5=0.52
00000000000001
Adding embeddings to index...
M=32, efConstruction=100, efSearch=50, metric=1: query_time=0.000507116317749023
4, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=100, efSearch=50, metric=2: query_time=0.005021810531616211,
precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=100, efSearch=100, metric=1: query_time=0.00100541114807128
9, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=100, efSearch=100, metric=2: query_time=0.000504255294799804
7, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=200, efSearch=10, metric=1: query_time=0.0, precision@5=0.52
00000000000001
Adding embeddings to index...
M=32, efConstruction=200, efSearch=10, metric=2: query_time=0.000998735427856445
3, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=200, efSearch=50, metric=1: query_time=0.002018690109252929
7, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=200, efSearch=50, metric=2: query_time=0.000506401062011718
8, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=200, efSearch=100, metric=1: query_time=0.001005172729492187
5, precision@5=0.5200000000000001
Adding embeddings to index...
M=32, efConstruction=200, efSearch=100, metric=2: query_time=0.002007484436035156
2, precision@5=0.5200000000000001

    == Best Query Time Results ==
1. Time: 0.0000s, M=8, efC=40, efS=10, Metric=L2
2. Time: 0.0000s, M=8, efC=40, efS=10, Metric=L1
3. Time: 0.0000s, M=8, efC=40, efS=50, Metric=L2
4. Time: 0.0000s, M=8, efC=100, efS=10, Metric=L2
5. Time: 0.0000s, M=8, efC=100, efS=10, Metric=L1
1. Precision: 0.5200, M=8, efC=40, efS=50, Metric=L1
2. Precision: 0.5200, M=8, efC=40, efS=100, Metric=L2
3. Precision: 0.5200, M=8, efC=40, efS=100, Metric=L1
4. Precision: 0.5200, M=8, efC=100, efS=50, Metric=L2
5. Precision: 0.5200, M=8, efC=100, efS=50, Metric=L1
Best Precision@5 [(8, 40, 50, 2, 0.0010035037994384766, 0.5200000000000001), (8,
40, 100, 1, 0.0005042552947998047, 0.5200000000000001), (8, 40, 100, 2, 0.0010046
958923339844, 0.5200000000000001), (8, 100, 50, 1, 0.0, 0.5200000000000001), (8,
100, 50, 2, 0.0, 0.5200000000000001)]
```

Task 5: Search for Dogs that look like you.

Use an image of yourself, and find dogs that look like you.

Alternatively: Search the internet and find at least 5 images of people, who look similar to dogs. For each image,

- Transform the image using CLIP
- Perform a search using any faiss index and plot the results.

```
In [24]: def dogs_that_look_like_you(embeddings, k):
    # Answer:
    query_path = "queries_task5"
    query_embeddings, query_file_names = create_query_embeddings(query_path)
    hswg_index, hswg_indices = index_and_query_hswg(embeddings, query_embeddings)
    plot_results(query_file_names, hswg_indices, k)

return None

dogs_that_look_like_you(embeddings, k)
```

The code defines a function `dogs_that_look_like_you` that takes embeddings and a k-value as input. It uses a pre-defined path for queries, creates query embeddings, indexes and queries the HSWG index, and plots the results. The function returns None. The function is then called with the embeddings and k.

The output shows four rows of search results:

- Row 1:** Query is a cartoon dog (Weimaraner). Results: 1-NN: Weimaraner (actual Weimaraner), 2-NN: Labrador_retriever (yellow Labrador), 3-NN: Italian_greyhound (Italian greyhound), 4-NN: redbone (redbone hound), 5-NN: basset (basset hound).
- Row 2:** Query is a cartoon man (Afghan hound). Results: 1-NN: Afghan_hound (actual Afghan hound), 2-NN: Afghan_hound (long-haired Afghan hound), 3-NN: black-and-tan_coonhound (black and tan coonhound), 4-NN: basset (basset hound), 5-NN: golden_retriever (golden retriever).
- Row 3:** Query is a cartoon woman (Italian greyhound). Results: 1-NN: Italian_greyhound (actual Italian greyhound), 2-NN: basset (basset hound), 3-NN: Weimaraner (actual Weimaraner), 4-NN: redbone (redbone hound), 5-NN: basset (basset hound).
- Row 4:** Query is a cartoon woman (Basset hound). Results: 1-NN: Italian_greyhound (actual Italian greyhound), 2-NN: Weimaraner (actual Weimaraner), 3-NN: Japanese_spaniel (Japanese spaniel), 4-NN: basset (basset hound), 5-NN: basset (basset hound).



Task 6: Text-To-Image Retrieval

We will now use CLIP text embeddings to find the top-3 matching images.

You are given a set of sentences. Your task is to:

- Generate CLIP-embedddings for each sentence
- Search for the Top-3 images using `IndexFlatL2` index and Metric `METRIC_L1`

Please refer to [1] for additional information on how to generate text-embeddings.

[1] <https://github.com/openai/CLIP>

You may add additional sentences to do further experiments.

```
In [30]: k = 3 # Number of neighbors
sentences = ["a lion",
             "a dog and a girl",
             "a dog and a man",
             "a toy and a dog",
             "a group of people"]

def text_to_image(embeddings, sentences, k):
    metric = faiss.METRIC_L1

    # Answer:
    index_flat = None
    indices = None # used to store the result

    # CLIP
    text_tokens = clip.tokenize(sentences).to(device)
    with torch.no_grad():
        text_embeddings = model.encode_text(text_tokens)
    text_embeddings = text_embeddings.cpu().numpy().astype('float32')

    index_flat = faiss.IndexFlatL2(embeddings.shape[1])
    # IndexFlatL2 is not suitable for L1 distance
    # Pre process the embeddings to simulate L1 distance
```

```

# Both embeddings
processed_embeddings = np.sign(embeddings) * np.sqrt(np.abs(embeddings))
processed_text_embeddings = np.sign(text_embeddings) * np.sqrt(np.abs(text_e

# Use processed embeddings
index_flat.add(processed_embeddings)
# Use provessed embeddings
distances, indices = index_flat.search(processed_text_embeddings, k)

return indices

indices = text_to_image(embeddings, sentences, k)

if flat_indices is None:
    print("You must implement Task 6")
else:
    assert precision_at_k(flat_indices, 1) == 1.0
    plot_results(sentences, indices, k)

```

a lion



a dog and a girl



a dog and a man



a toy and a dog



a group of people



Hand in your solution via Moodle

Submit two files:

- This notebook
- A html-export of this notebook